

BURNING DOWN THE HOUSE: KEEPING YOUR POSTGRESQL DATA SAFE



YOU LOVE YOUR DATA

**WHICH IS WHY YOU STORE
IT IN A DATABASE**

BUT...

**WHAT HAPPENS IF
YOUR HARDWARE FAILS?**



**WHAT HAPPENS IF
YOUR SOFTWARE FAILS?**

**“AT 9:45 EST A USER TRIGGERED AN UNSCOPED
DELETION OF ALL HISTORICAL PERIOD RECORDS [...]”**

DEAD MAN'S SNITCH OUTAGE POST MORTEM

**NOT IF
BUT WHEN**

**CHOOSE A
BACKUP STRATEGY**

PAY ATTENTION TO...

**HOW MUCH DATA CAN
YOU AFFORD TO LOSE?**

**HOW MUCH DOWN TIME
CAN YOU AFFORD?**

WAYS TO BACK UP YOUR DATABASE

HAVE SOMEONE ELSE DO IT

✓ IT'S SOMEONE ELSE'S
PROBLEM NOW

✗ LESS CONTROL

✗ CAN COST A LOT

✗ LESS FLEXIBILITY

FILE-BASED BACKUP METHODS

✓ EASY

✓ WIDELY AVAILABLE

TOOLS

✗ REQUIRES POSTGRES
SQL SERVER TO BE SHUT DOWN


```
$ pg_ctl stop
```

```
$ rsync $PGDATA /path/to/backup
```

```
$ pg_ctl start
```

pg_dump[all]

✓ CREATES LOGICAL
BACKUP

✓ FLEXIBLE

✓ NO DOWN TIME REQUIRED

✓ MINIMAL IMPACT

✗ IMPRACTICAL FOR
LARGE DATABASES

```
$ pg_dump \  
--format=custom \  
--exclude-table-data=stats \  
--compress=9 \  
--jobs=4 \  
${DB_NAME}.pg_dump  
$ # Repeat for all Databases
```

```
$ pg_restore \  
--jobs=4 \  
--dbname=${DB_NAME} \  
/path/to/backup.pg_dump  
$ # Get a cup of coffee
```


HOT STANDBY

- ✓ CONTINUOUS
- ✓ USE THEM TO DISTRIBUTE READ LOAD
- ✓ GOOD FOR FAST FAILOVER

- ✗ ONLY PROTECTS AGAINST HARDWARE FAILURE
- ✗ REQUIRES AT LEAST ONE MORE SERVER

```
# In postgresql.conf:  
wal_level = hot_standby
```

```
# Set these to something > 0  
max_wal_senders = 5  
wal_keep_segments = 64
```

```
$ pg_basebackup \
--host=localhost \
--username=replication_role
--format=plain \
--xlog-method=stream \
--pgdata=${STANDBY_PGDATA} \
--progress
```

```
# postgresql.conf on standby:  
hot_standby = on  
hot_standby_feedback = on  
  
max_standby_streaming_delay = 10s
```



```
standby_mode = 'on'  
primary_conninfo = 'host=db-primary port=5432  
user=replication_role sslmode=require'  
trigger_file = '/usr/local/pgsql/data/  
primary.trigger'
```

```
$ pg_ctl start
```

```
LOG:  entering standby mode
```

```
... then some time later ...
```

```
LOG:  consistent recovery state reached
```

```
LOG:  database system is ready to  
accept read only connections
```

PITR-BASED BACKUPS

✓ CONTINUOUS BACKUP

✓ HIGH RECOVERABILITY

✓ CLUSTER-BASED

✗ COMPLEX

✗ INCREASED I/O

✗ NEEDS LOTS OF STORAGE

✗ ARCHITECTURE-
DEPENDANT

```
archive_command = 'cp %p /path/to/archive/%f'
```



```
$ pg_basebackup \
--host=localhost \
--username=replication_role
--format=plain \
--pgdata=${BACKUP_PGDATA} \
--progress
```

```
# Restore/untar latest base backup
```

```
# Create a recovery.conf:
```

```
restore_command = 'cp /path/to/archive/%f %p'
```

```
# Start PostgreSQL
```

BARMAN
OMNIPITR

3RD PARTY TOOLS

✓ REPLICATE BETWEEN

POSTGRESQL VERSIONS

✓ VERY FLEXIBLE

✓ PAID SUPPORT

✗ COMPLICATED

✗ SETUP & MAINTENANCE

✗ COSTS MONEY

RECOMMENDATIONS

DAILY pg_dump
HOT STANDBY

**STORE BACKUPS
SOMEWHERE ELSE**

TEST YOUR BACKUPS

SERIOUSLY, TEST THEM

POSTGRESQL.ORG/DOCS/CURRENT/

**BACK UP YOUR DATA
(NOT JUST YOUR DATABASE)**

THX!

NUCLEARSQUID.COM / @NUCLEARSQUID / GITHUB.COM/CYPHER